

Periodic Neighbor Search Architecture Manual

Periodic Geometry, Exact Cell Lists, Verlet Reuse, and Staged Implementation

mdstats

2026-07-15 (theoretical background revision)

Contents

Purpose and status	3
Theoretical background	3
Periodic configuration space	3
The cutoff graph is the scientific object	4
Spatial factorization by linked cells	4
Temporal coherence and the Verlet skin	4
Deforming cells: affine and nonaffine motion	4
Exactness, conservatism, and performance	5
Implementation status	6
Accepted scope and constraints	6
Fixed atom-population promise	6
Scientific contract remains unchanged	7
Dense backend is retained	7
Current baseline	7
Target subsystem architecture	8
Mathematical conventions	9
Cell and coordinate convention	9
Strict cutoff convention	9
Metric tensor	9
Algorithmic provenance and attribution boundaries	9
Exact periodic cell-list design	10
Reduced search basis	10
Fractional linked cells	10
Metric-aware bin stencil	11
Candidate generation	11
Duplicate control	11
Minimum-image uniqueness condition	11
Verlet candidate cache	12
Separation of roles	12
Candidate-list radius	12
Cache contents	12
Request signature	13
Fixed-cell validity criterion	13
Deformation-aware cache validity	13
Affine cell deformation	13
Nonaffine displacement	14
Accepted safety margin	14
Cache invalidation independent of the margin	14
Hysteretic and reference connectivity	15
Separate geometric and bond state	15
Single-pass hysteresis	15
Internal objects and module boundaries	15

Cache lifecycle	17
Staged implementation plan	18
Stage S0 - Dense oracle and baseline harness	18
Goal	18
Work	18
Acceptance gate	18
Deliverable	18
Stage S1 - Exact triclinic cell-list backend	18
Goal	18
Work	18
Acceptance gate	19
Deliverable	19
Implemented result	19
Stage S2 - Fixed-cell Verlet cache	19
Goal	19
Work	19
Acceptance gate	20
Deliverable	20
Implemented result	20
Stage S3 - Deformation-aware Verlet reuse	20
Goal	20
Work	20
Acceptance gate	20
Deliverable	21
Implemented result	21
Stage S4 - Integration, automatic policy, benchmarks, and specifications	21
Goal	21
Work	21
Acceptance gate	21
Deliverable	22
Implemented result	22
Test strategy	22
Dense equivalence is the primary oracle	22
Randomized property tests	22
Basis-invariance tests	22
Cache-completeness tests	23
Consumer-level regression	23
Benchmark strategy	23
Systems	23
Sizes	23
Metrics	23
Auto-backend policy	24
Diagnostics and provenance	24
Specification and documentation impact	24
New normative specification	24
Existing specifications	24
Risks and mitigations	25
Missing candidates in skewed cells	25
Wrong periodic image shift	25
Invalid cache reuse	25
Wrapped-coordinate displacement errors	25
Overly large candidate lists	25

Premature automatic selection	25
Scientific-module coupling	25
Explicitly deferred features	25
Definition of done	26
Completion status	26
Appendix A - Rebuild and reuse pseudocode	26
Appendix B - Stage gate summary	27
Appendix C - Accepted design decisions	27
References	28

Purpose and status

This document records the accepted high-level plan for replacing the current blocked dense neighbor search in `mdstats` with an exact periodic cell-list backend and a trajectory-aware Verlet candidate cache.

The plan is intentionally staged. Each stage must be independently testable and must preserve the existing scientific neighbor contract before the next stage is started. The document is a design and implementation reference, not yet a final module specification. Detailed function signatures may be refined during each stage, but the mathematical invariants, ownership boundaries, and acceptance gates in this document should remain stable unless an explicit revision is made.

The optimization is motivated by the current per-frame scaling

$$T_{\text{dense}} \sim \sum_{(A,B) \in \mathcal{C}} n_A n_B,$$

where $\mathcal{C}$ is the set of registered species pairs and n_A is the number of atoms of species A in the active scope. Species pruning and NumPy vectorization reduce the constant factor, but the worst-case work remains quadratic in system size.

The target architecture should approach linear work at fixed density and finite cutoff:

$$T_{\text{target}} \sim O(N z_{\text{list}}),$$

where z_{list} is the average number of cached candidates per atom inside the physical cutoff plus skin.

Theoretical background

Periodic configuration space

A periodic cell with row-vector matrix H generates the lattice

$$\mathcal{L} = \{\mathbf{n}H : \mathbf{n} \in \mathbb{Z}^3\}.$$

Cartesian points that differ by a lattice vector represent the same physical point under periodic boundary conditions. Fractional coordinates therefore live on the quotient space

$$\mathbb{R}^3 / \mathbb{Z}^3,$$

which is a three-dimensional torus when all axes are periodic. A pair displacement is not one vector but the equivalence class

$$\Delta \mathbf{r}(\mathbf{n}) = (\mathbf{s}_j - \mathbf{s}_i + \mathbf{n})H.$$

The minimum-image distance is the closest-vector problem

$$r_{ij} = \min_{\mathbf{n} \in \mathbb{Z}^3} \| (\mathbf{s}_j - \mathbf{s}_i + \mathbf{n})H \|.$$

For skewed cells, componentwise rounding of fractional coordinates is not a general solution to this minimization. The cell metric

$$G = HH^T$$

encodes the full geometry through

$$\|\Delta \mathbf{r}\|^2 = \Delta \mathbf{s} G \Delta \mathbf{s}^T.$$

The cutoff graph is the scientific object

For a physical cutoff r_c , the neighbor relation defines a graph

$$E(r_c) = \{(i, j, \mathbf{m}_{ij}) : r_{ij} < r_c\}.$$

Dense search, cell lists, and Verlet caches are alternative algorithms for constructing the same graph. Acceleration is scientifically valid only when it preserves every accepted pair, distance, displacement vector, and periodic image shift. The dense $O(N^2)$ calculation is therefore retained as an oracle even after faster backends are introduced.

At fixed number density and finite cutoff, the expected number of physical neighbors per atom remains bounded. The computational target is consequently

$$O(Nz),$$

where z is the mean local candidate count, rather than testing all $N(N-1)/2$ pairs.

Spatial factorization by linked cells

The linked-cell method partitions space into bins and searches only bin pairs that can approach within the list radius. Its classical lineage begins with Quentrec and Brot [1]. In an orthogonal box, nearby bins can be enumerated by a simple Cartesian stencil. In a triclinic cell, fractional bin offsets must be judged using the metric G .

For a pair of bins, let $B \subset \mathbb{R}^3$ be the box of possible fractional differences between points in those bins. The exact question is

$$\min_{\Delta \mathbf{s} \in B} \Delta \mathbf{s} G \Delta \mathbf{s}^T < R^2.$$

If the lower bound is at least R^2 , the bin pair cannot contain a candidate and may be skipped. Otherwise it must be searched. This is the mathematical role of the metric-aware stencil: it proves exclusion rather than guessing a fixed set of neighboring bin indices.

A reduced lattice basis can make the search representation more compact, but the physical result must be reconstructed in the original cell basis. Basis reduction changes efficiency, not the periodic structure being analyzed.

Temporal coherence and the Verlet skin

Molecular trajectories are continuous in time, so the neighbor graph usually changes much more slowly than the coordinates. The Verlet method [2] exploits this by building a candidate list at the enlarged radius

$$R = r_c + r_{\text{skin}}.$$

At later frames, exact current distances are recomputed only for cached candidates and filtered by the physical cutoff r_c . A pair absent from the candidate list was separated by at least R at rebuild time. In a fixed cell, if every atom has moved by at most d_{max} , two omitted atoms can approach by at most $2d_{\text{max}}$. Completeness is therefore guaranteed while

$$2d_{\text{max}} < r_{\text{skin}},$$

up to the explicit numerical tolerance. Automatic displacement-triggered list updating belongs to the established Verlet-list literature [3].

The cache is not a stored bond graph. It is a conservative superset of possible neighbors, and every reused frame still undergoes an exact minimum-image cutoff test.

Deforming cells: affine and nonaffine motion

For a rebuild cell H_0 and current cell H_t , the row-vector deformation map is

$$F_t = H_0^{-1} H_t.$$

An atom position can be decomposed as

$$\mathbf{r}_i(t) = \mathbf{r}_i(0)F_t + \mathbf{u}_i(t),$$

where the first term is affine cell deformation and $\mathbf{u}_i$ is nonaffine atomic motion. The smallest singular value supplies a global contraction bound,

$$\| \mathbf{x} F_t \| \geq \sigma_{\min}(F_t) \| \mathbf{x} \| .$$

For a pair type A, B , an omitted rebuild pair at distance at least $r_{AB} + r_{\text{skin}}$ remains outside the physical cutoff if

$$\sigma_{\min}(F_t)(r_{AB} + r_{\text{skin}}) - u_A^{\max}(t) - u_B^{\max}(t) > r_{AB}.$$

This is the basis of the implemented deformation-aware margin. The use of a global singular-value contraction bound together with species-resolved nonaffine displacement maxima is an `mdstats` derivation. Published work on general and deforming periodic neighbor searches [4-7] supplies important prior context but not this exact theorem.

Exactness, conservatism, and performance

The subsystem follows three distinct principles:

- **exactness:** no valid physical neighbor may be omitted;
- **conservatism:** uncertain bin pairs or cache states fall back to more work, not fewer candidates;
- **performance:** reduced bases, spatial bins, and temporal reuse lower cost only after exactness is established.

This hierarchy explains why the architecture retains dense comparison tests, explicit fallback events, and complete provenance. A faster candidate generator is acceptable only when the final neighbor graph is indistinguishable from the dense definition.

Implementation status

Stage	Status	Package	Notes
S0	Complete	0.14.0a0	Dense backend made explicit; immutable results, canonical comparison utilities, randomized scalar-oracle tests, and baseline benchmarks added.
S1	Complete	0.14.0a1	Exact single-frame triclinic cell-list backend; optional reduced search basis, perpendicular-height bins, exact metric stencil, original-basis MIC output, and dense-equivalence benchmark.
S2	Complete	0.14.0a2	Request-keyed fixed-cell Verlet candidate cache, exact current MIC reuse, cache statistics, and one-pass nested connectivity thresholds.
S3	Complete	0.14.0a3	Explicit singular-value and species-aware nonaffine deformation bound, rigid-rotation reuse, condition-number validation, and variable-cell equivalence tests.
S4	Complete	0.14.1	Public policy integration, deterministic automatic selection, semantics-aware cache resolution, repeated-zero-reuse shutoff, exact fallback, unified diagnostics, representative benchmarks, and release validation.

Stages S0-S4 are complete. The exact dense backend remains the scientific oracle, while the public `NeighborSearchOptions` policy selects dense or cell-list execution and enables request-keyed Verlet reuse for multi-frame cell-list workloads. RDF, coordination, bond angle, and every distance-based connectivity mode share this policy and emit the same diagnostic schema. Explicit dense and cell-list overrides remain available, and every automatic fallback is exact and recorded.

Accepted scope and constraints

Fixed atom-population promise

`AtomisticFrameCollection` provides the fixed-population promise required by the cache:

- every frame contains the same number of atoms;
- canonical atom indices are stable;

- persistent source IDs are stable when available;
- species and masses are stable at each canonical atom index;
- periodic-boundary flags are stable within the collection;
- coordinates and cells may vary from frame to frame;
- atomic connectivity and framework topology may change.

The neighbor subsystem may therefore cache candidate pairs by canonical atom index. It must not assume that geometry, connectivity, or phase remains fixed.

Scientific contract remains unchanged

Scientific consumers specify **what** neighborhood they need. The neighbor subsystem decides **how** to obtain it.

Consumers include:

- pair RDF;
- coordination distributions;
- bond-angle distributions;
- atomic connectivity;
- future local-environment and topology modules.

These modules must not own:

- cell-list construction;
- lattice reduction;
- skin-distance validity tests;
- cache rebuild decisions;
- request-signature hashing;
- reduced-basis image-shift conversion.

They continue to consume the same conceptual output:

canonical atom indices
 minimum-image vectors
 minimum-image distances
 periodic image shifts
 CSR or equivalent pair grouping

Dense backend is retained

The existing blocked dense search remains permanently available as:

- a correctness oracle;
- an efficient backend for small systems;
- a fallback for unsupported or pathological geometries;
- a regression reference for every optimized backend.

The optimized implementation must never remove the ability to force the dense backend explicitly.

Current baseline

The current private neighbor kernel performs the following work for each frame:

1. resolve active atom selections and species-pair cutoffs;
2. loop over registered species pairs;
3. form dense center-candidate displacement blocks;
4. apply periodic minimum-image geometry vectorially;
5. filter distances by the strict physical cutoff;
6. return a compact CSR-like result.

Its strengths are:

- deterministic behavior;

- exact frame-local minimum-image geometry;
- orthogonal and triclinic support through the existing MIC operation;
- species and scope pruning;
- bounded temporary memory through center blocking;
- compact accepted-neighbor storage.

Its limitations are:

- dense candidate generation for every registered species pair;
- no spatial cell list;
- no candidate reuse between frames;
- two separate geometric scans in the current hysteretic and reference paths;
- no diagnostics for candidate efficiency or rebuild frequency.

The new design must preserve the strengths while replacing candidate generation and trajectory reuse.

Target subsystem architecture




Figure 1: Target neighbor-search architecture

The intended ownership boundary is:

```

scientific module
|
v
normalized neighbor request
|
v
persistent NeighborSearchSession
|

```



```

+--> compatible valid cache: evaluate cached pairs
|
`--> no valid cache: rebuild with exact cell list
|
v
apply physical cutoff or connectivity rule
|
v
unchanged neighbor result contract

```

The session owns one or more caches keyed by normalized request signature. The initial implementation should use one cache per exact request signature rather than attempting cross-analysis superset sharing.

Mathematical conventions

Cell and coordinate convention

The three lattice vectors are rows of the cell matrix

$$H = \begin{pmatrix} \mathbf{a}^T \\ \mathbf{b}^T \\ \mathbf{c}^T \end{pmatrix}.$$

Fractional row vectors map to Cartesian coordinates by

$$\mathbf{r} = \mathbf{s}H.$$

A periodic image displacement from atom i to atom j is

$$\mathbf{d}_{ij} = (\mathbf{s}_j - \mathbf{s}_i + \mathbf{m}_{ij})H,$$

where $\mathbf{m}_{ij} \in \mathbb{Z}^3$ is zero along nonperiodic axes.

All final accepted vectors, distances, and image shifts must be derived from the same exact minimum-image operation.

Strict cutoff convention

A physical neighbor satisfies

$$r_{ij} < r_{AB},$$

not $r_{ij} \leq r_{AB}$. The strict relation is part of the existing neighbor contract and must remain unchanged.

Metric tensor

The fractional-space metric is

$$G = HH^T,$$

so that

$$\|\Delta \mathbf{r}\|^2 = \Delta \mathbf{s} G \Delta \mathbf{s}^T.$$

The off-diagonal elements of G encode cell skew. They are the reason that a hard-coded nearest-bin stencil is unsafe for general triclinic cells.

Algorithmic provenance and attribution boundaries

The architecture combines several established algorithmic lineages:

- Quentrec and Brot's cell-linked-list decomposition [1];
- Verlet buffering [2] and automatic displacement-triggered updating [3];
- arbitrary-periodic-box and parallelepiped pair-list geometry [4-6];
- dynamically deforming periodic cell-list methods [7];

- ASE’s lattice software [8] and the Nguyen-Stehlé low-dimensional reduction algorithm used by its Minkowski-reduction implementation [9].

These references establish foundations and neighboring prior art. The exact metric-box active-set stencil, original-basis image reconstruction, immutable request-keyed cache model, and species-resolved smallest-singular-value S3 margin are mdstats-specific design results. In particular, reference [7] is not claimed as the source of the S3 cache-validity theorem.

Exact periodic cell-list design

Reduced search basis

A highly skewed crystallographic basis can make an otherwise correct linked-cell stencil unnecessarily large. The search backend should therefore construct an equivalent reduced lattice basis. The production implementation requests this basis from ASE [8], whose low-dimensional reduction routine follows Nguyen and Stehlé [9].

$$H_s = UH,$$

where U is an integer unimodular matrix:

$$U \in \text{GL}(3, \mathbb{Z}), \quad \det U = \pm 1.$$

The physical lattice is unchanged. Only the internal search representation is altered.

With the row-vector convention, search-basis fractional coordinates are

$$\mathbf{s}_s = \mathbf{s}U^{-1}.$$

If an image shift $\mathbf{m}_s$ is found in the search basis, the reported shift in the original cell basis is

$$\mathbf{m} = \mathbf{m}_s U.$$

The original cell remains authoritative for all public outputs. Lattice reduction is an internal acceleration device, not a change of scientific coordinates.

The first implementation may use a well-tested lattice-reduction routine from a trusted numerical dependency. If no dependable routine is available for mixed PBC, the backend may skip reduction and rely on the metric stencil; correctness must not depend on reduction.

Fractional linked cells

Atoms are wrapped and binned in the search-basis fractional domain. For fully periodic axes this domain is $[0, 1)^3$. For nonperiodic axes, finite bin ranges are constructed from the active Cartesian or fractional extent with explicit boundary padding.

The bin-count heuristic should use perpendicular lattice heights rather than basis-vector lengths. For

$$V = |\mathbf{a} \cdot (\mathbf{b} \times \mathbf{c})|,$$

define

$$h_a = \frac{V}{\|\mathbf{b} \times \mathbf{c}\|}, \quad h_b = \frac{V}{\|\mathbf{c} \times \mathbf{a}\|}, \quad h_c = \frac{V}{\|\mathbf{a} \times \mathbf{b}\|}.$$

A practical bin-count heuristic is

$$n_\alpha = \max\left(1, \left\lfloor \frac{h_\alpha}{\ell_{\text{target}}} \right\rfloor\right),$$

where ℓ_{target} is of order the maximum list radius. Correctness must not depend on this exact heuristic because the metric-aware stencil and exact final distance test remain authoritative.

Metric-aware bin stencil

For a bin-offset vector $\mathbf{k}$, let $D_{\mathbf{k}}$ be the rectangular set of all possible fractional displacement vectors between points in the two bins. The bin pair must be searched whenever

$$\min_{\Delta \mathbf{s} \in D_{\mathbf{k}}} \Delta \mathbf{s} G_{\mathbf{s}} \Delta \mathbf{s}^T < R_{\max}^2,$$

where

$$G_{\mathbf{s}} = H_{\mathbf{s}} H_{\mathbf{s}}^T$$

and

$$R_{\max} = \max_{(A,B)} R_{AB}.$$

The list radius for pair type A, B is

$$R_{AB} = r_{AB} + r_{\text{skin}}.$$

The three-dimensional box-constrained quadratic minimization is small and deterministic. A suitable exact implementation can enumerate the 3^3 active sets in which each coordinate is free, fixed to its lower bound, or fixed to its upper bound. For each active set:

1. solve the free-coordinate quadratic optimum;
2. reject infeasible solutions;
3. evaluate the metric distance;
4. retain the smallest feasible value.

The stencil is geometry-dependent but atom-independent. It is built once per cell-list rebuild and may be cached with the bin layout.

Candidate generation

For each occupied bin:

1. visit only bins in the metric-derived stencil;
2. respect periodic wrapping only along periodic axes;
3. generate canonical unordered atom pairs;
4. reject unregistered species pairs;
5. apply the pair-specific list radius R_{AB} ;
6. compute exact minimum-image vectors and original-basis image shifts;
7. store candidate atom pairs inside the list radius.

The stencil may generate false candidates; the exact pair-specific list-radius test removes them. It must never generate a false negative.

Duplicate control

Candidate identity is initially the canonical atom-index pair

$$(i, j), \quad i < j.$$

Under the initial unique-minimum-image regime, each atom pair contributes at most one relevant periodic image. Duplicate discoveries from multiple bin paths must be removed deterministically.

Minimum-image uniqueness condition

As of 0.19.73a0, the production neighbor subsystem uses the exact unique-image radius rather than the historical perpendicular-face-height bound. For each frame, define the periodic translation lattice

$$\Lambda_{\text{PBC}}(H) = \{\mathbf{n}H : \mathbf{n} \in \mathbb{Z}^3, n_{\alpha} = 0 \text{ on nonperiodic axes}\}.$$

The **largest list radius**, not only the largest physical cutoff, must satisfy

$$R_{\max} \leq \frac{1}{2} \min_{\ell \in \Lambda_{\text{PBC}}(H) \setminus \{0\}} \|\ell\|.$$

The pair inclusion rule remains strict, `distance < cutoff`, so a list radius exactly equal to the bound does not retain an ambiguous boundary pair. Across selected frames, `mdstats` uses the minimum frame-local radius.

The shortest translation is obtained from ASE’s low-dimensional Minkowski reduction [8,9]. `mdstats` validates the returned integer unimodular transform, its consistency with the reduced cell, preservation of the periodic sublattice, and ASE’s reduced-basis certificate. In a Minkowski-reduced basis, the shortest periodic basis vector is a shortest lattice vector.

Perpendicular cell heights remain part of the cell-list binning and metric- stencil construction, but they no longer limit the scientific neighbor cutoff. For skewed primitive cells this removes an unnecessary rejection. For example, a near-60° LTA primitive cell with approximately 17.36 Å basis-vector lengths has an exact unique-image radius near 8.68 Å, while the old face-height bound was only about 7.09 Å.

The safety condition ensures that one minimum-image representation per canonical atom pair is sufficient. If the list radius violates the supported bound, the implementation should request a smaller skin, fall back only when the existing backend has identical documented semantics, or raise an explicit unsupported- regime error. Silent omission of additional periodic images is not acceptable.

Verlet candidate cache

Separation of roles

The cell list and Verlet cache perform different tasks:

- the cell list builds a sparse candidate set efficiently;
- the Verlet cache reuses that candidate set across nearby frames.

The cell list is invoked only when a cache is created or rebuilt. On reuse frames, exact geometry is evaluated only for cached candidate atom pairs.

Candidate-list radius

For each physical pair cutoff r_{AB} and one global initial skin $r_{\text{skin}} > 0$,

$$R_{AB} = r_{AB} + r_{\text{skin}}.$$

This is the classical buffered candidate radius of the Verlet-list method [2].

The cache stores all registered atom pairs satisfying

$$r_{ij}(t_0) < R_{AB}$$

at the rebuild frame t_0 .

The physical result at every frame still uses r_{AB} , not R_{AB} . The skin is only a completeness margin.

Cache contents

A cache should contain, at minimum:

```
@dataclass(frozen=True, slots=True)
class VerletPairCache:
    request_digest: str
    reference_frame_index: int
    reference_positions: FloatArray      # (N, 3)
    reference_fractional_positions: FloatArray
    reference_cell: FloatArray           # (3, 3)
    candidate_atom_i: IntArray           # (P,)
    candidate_atom_j: IntArray           # (P,)
    candidate_pair_type_ids: IntArray    # (P,)
    physical_pair_radii: FloatArray
    list_pair_radii: FloatArray
```

```
skin: float
diagnostics: NeighborCacheStatistics
```

The exact field layout may change, but the cache must retain enough information to:

- prove geometric validity;
- evaluate current exact distances;
- invalidate on request changes;
- report rebuild and efficiency statistics.

A cached image shift may be stored as a hint, but it must not be treated as permanently authoritative. The current minimum-image vector and image shift are recomputed on every reuse frame.

Request signature

A cache is reusable only for the same normalized scientific request. The request digest should include:

- canonical atom selection or connectivity scope;
- center and candidate semantics;
- registered species-pair radii;
- PBC flags;
- skin distance;
- backend and cell-list options that affect completeness;
- atom identities and species schema;
- canonical schema version.

A different request signature causes automatic rebuild or creation of a separate cache. Scientific consumers never manually decide whether a cache is compatible.

Fixed-cell validity criterion

At a rebuild frame, any omitted pair satisfies

$$r_{ij}(t_0) \geq r_{AB} + r_{\text{skin}}.$$

If each atom has moved no more than d_{max} , the pair distance can shrink by at most $2d_{\text{max}}$. Therefore, the cache remains complete while

$$2d_{\text{max}} < r_{\text{skin}}.$$

The implementation should rebuild when

$$2d_{\text{max}} \geq r_{\text{skin}} - \varepsilon,$$

where ε is a documented numerical safety tolerance.

Wrapped positions must not be subtracted directly. Displacement must be derived from unwrapped coordinates or reference-relative fractional unwrapping.

Deformation-aware cache validity

Affine cell deformation

Let the rebuild cell be H_0 and the current cell be H_t . Define the row-vector affine deformation map

$$F_t = H_0^{-1} H_t.$$

A reference Cartesian pair vector transforms affinely as

$$\mathbf{x}_{\text{aff}}(t) = \mathbf{x}_0 F_t.$$

Let

$$\sigma_{\min}(F_t)$$

be the smallest singular value of F_t . Then every vector satisfies

$$\|\mathbf{x}_0 F_t\| \geq \sigma_{\min}(F_t) \|\mathbf{x}_0\|.$$

The singular-value bound handles isotropic strain, orthorhombic strain, shear, and rigid cell rotation in one expression.

Nonaffine displacement

Write the current atom position as

$$\mathbf{r}_i(t) = \mathbf{r}_i(0)F_t + \mathbf{u}_i(t),$$

where $\mathbf{u}_i$ is the nonaffine displacement from the affinely deformed reference configuration.

Using fractional coordinates,

$$\mathbf{u}_i(t) = \Delta \mathbf{s}_i(t) H_t,$$

where $\Delta \mathbf{s}_i(t)$ is the correctly unwrapped fractional displacement from the rebuild frame.

For each species A , define

$$u_A^{\max}(t) = \max_{i \in A} \|\mathbf{u}_i(t)\|.$$

The species-aware bound avoids allowing highly mobile spectator ions to consume the safety margin of a relatively immobile framework pair type.

Accepted safety margin

For pair type A, B , define

$$M_{AB}(t) = \sigma_{\min}(F_t)(r_{AB} + r_{\text{skin}}) - r_{AB} - u_A^{\max}(t) - u_B^{\max}(t).$$

The cache is valid when

$$\min_{(A,B) \in \mathcal{C}} M_{AB}(t) > \varepsilon$$

for all active registered pair types.

When $H_t = H_0$,

$$\sigma_{\min} = 1,$$

and the condition reduces to the ordinary fixed-cell criterion.

This policy permits safe reuse under modest anisotropic strain and shear. It is conservative by construction and requires only one 3×3 singular-value decomposition plus species-wise displacement maxima per frame.

Cache invalidation independent of the margin

The cache must be rebuilt or discarded if any of the following occurs:

- atom count or canonical atom identities change;
- species assignments change;
- PBC flags change;
- active scope or atom selection changes;
- cutoff registry changes;
- skin or backend options change;
- the current cell is singular or numerically ill-conditioned;
- fractional displacement unwrapping is ambiguous;
- request schema or cache schema changes;
- the safety margin is not finite or is below tolerance.

A change in the internally selected reduced basis does not by itself invalidate a valid atom-pair cache. The reduced basis is needed only at rebuild time.

Hysteretic and reference connectivity

Separate geometric and bond state

The Verlet cache owns geometric candidate completeness:

Which atom pairs could possibly matter under the largest physical threshold?

Hysteretic connectivity owns previous bond state:

Which candidate pairs were bonded in the preceding frame?

These states must not be conflated.

Single-pass hysteresis

For pair type A, B with

$$r_{AB}^{\text{form}} < r_{AB}^{\text{break}},$$

the list radius is

$$R_{AB} = r_{AB}^{\text{break}} + r_{\text{skin}}.$$

One exact distance pass over cached candidates supports both decisions:

$$\text{previously bonded} \implies r < r_{AB}^{\text{break}},$$

$$\text{previously absent} \implies r < r_{AB}^{\text{form}}.$$

This replaces the current redundant formation and breaking scans.

Reference-retention and reference-formation rules should use the same one-pass principle, with the largest relevant threshold controlling the candidate cache.

Internal objects and module boundaries

The exact API will be specified during implementation. The intended internal separation is:

mdstats/analysis/_neighbors.py
stable facade and dense reference backend

mdstats/analysis/_cell_list.py
reduced-basis preparation
fractional bins
metric stencil
exact candidate construction

mdstats/analysis/_verlet_cache.py
candidate cache
request signature
fixed/deforming-cell validity
rebuild statistics

mdstats/analysis/_neighbor_engine.py
persistent search session
backend selection
request-keyed cache controller

Likely internal value objects are:

NeighborSearchRequest
NeighborSearchOptions
CellListLayout
CellListStencil
VerletPairCache
NeighborCacheStatistics
NeighborSearchSession

Scientific modules should require no knowledge of these internal decomposition details. At most, their trajectory loops may create and retain a shared search session instead of repeatedly calling a stateless helper.

Cache lifecycle




Figure 2: Verlet-cache lifecycle

Conceptual pseudocode:

```
function evaluate(frame, request):  
    key = normalize_and_digest(request)  
    cache = caches.get(key)  
  
    if cache is absent or not structurally compatible:  
        cache = rebuild_with_cell_list(frame, request)  
  
    else:
```

```

    margins = deformation_aware_margins(frame, cache, request)
    if any margin <= tolerance:
        cache = rebuild_with_cell_list(frame, request)

    geometry = evaluate_exact_MIC_for_cached_pairs(frame, cache)
    result = apply_physical_rules(geometry, request)
    update_statistics(cache, result)
    return result

```

Staged implementation plan




Figure 3: Staged implementation sequence

Each stage ends with a release-quality artifact and an explicit gate. A later stage must not begin until the preceding gate passes.

Stage S0 - Dense oracle and baseline harness

Goal

Harden the current dense implementation as the authoritative comparison backend and establish reproducible correctness and performance baselines.

Work

- formalize an explicit dense backend option;
- normalize pair-result sorting and comparison utilities;
- add randomized geometry generators;
- add a neighbor-result equality assertion that compares:
 - canonical pair identity;
 - image shift;
 - vector;
 - distance;
 - CSR grouping;
- record current timing and peak-memory baselines;
- document the current complexity and limits;
- add fixtures for orthogonal, triclinic, mixed-PBC, and boundary cases.

Acceptance gate

- all existing tests pass unchanged;
- dense output is deterministic under repeated runs;
- randomized dense tests are reproducible from stored seeds;
- the baseline report records system size, species counts, cutoff registry, runtime, and accepted-pair count;
- no scientific output changes.

Deliverable

A baseline package and short audit report. No optimized behavior is enabled yet.

Stage S1 - Exact triclinic cell-list backend

Goal

Implement an exact cell-list candidate generator that matches dense search on a single frame.

Work

- implement search-basis preparation and optional lattice reduction;
- implement fractional bin assignment;

- implement perpendicular-height bin heuristics;
- implement exact metric-aware bin stencil construction;
- implement periodic and mixed-PBC bin traversal;
- implement species-pair list-radius filtering;
- convert image shifts back to the original basis;
- deterministically deduplicate candidates;
- expose explicit backend selection:

```
neighbor_backend="dense"
neighbor_backend="cell_list"
```

- retain dense fallback for small or unsupported cases;
- do not implement trajectory caching yet.

Acceptance gate

For every test frame, sorted cell-list results must equal dense results exactly in pair identity and image shift and within numerical tolerance in vectors and distances.

Required geometries:

- cubic and orthorhombic cells;
- moderately skewed triclinic cells;
- deliberately highly skewed but nonsingular cells;
- atoms on or near periodic boundaries;
- mixed PBC;
- nonperiodic selections;
- multiple pair-specific cutoffs;
- identical-species and cross-species pairs;
- random atom permutations with stable canonical indices;
- small systems including zero and one accepted pair;
- dense local clusters;
- cutoff values near the unique-image limit.

No false negative is permitted. Extra bin candidates are allowed internally but must be removed by the exact list-radius test.

Deliverable

A package with the cell-list backend available only by explicit request, plus an updated internal-neighbor specification and dense-equivalence report.

Implemented result

Completed in 0.14.0a1. The implementation uses optional ASE Minkowski reduction, perpendicular-height fractional bins, an exact active-set metric-box stencil, deterministic candidate deduplication, and authoritative original-cell MIC evaluation. The complete normative specification is docs/specs/analysis/_cell_list_spec.md.

Stage S2 - Fixed-cell Verlet cache

Goal

Reuse cell-list candidates across fixed-cell trajectory frames.

Work

- implement normalized request signatures;
- implement VerletPairCache;
- implement persistent NeighborSearchSession;
- implement fixed-cell displacement tracking;
- rebuild when

$$2d_{\max} \geq r_{\text{skin}} - \varepsilon;$$

- recompute exact current MIC geometry for cached pairs;
- expose opt-in cache configuration;
- report rebuild and reuse statistics;
- merge hysteretic formation/breaking scans into one candidate pass;
- keep scientific hysteresis state outside the cache.

Acceptance gate

For every frame of every fixed-cell test trajectory:

$$\text{cached result} = \text{fresh cell-list result} = \text{fresh dense result} .$$

Required trajectories:

- harmonic solid motion;
- diffusive liquid motion;
- periodic boundary crossings;
- a pair that enters the physical cutoff just before rebuild;
- a pair omitted from the cache that approaches the cutoff without violating the skin bound;
- exact rebuild-threshold cases;
- hysteretic bond formation and breaking;
- request changes that force invalidation;
- noncontiguous frame requests evaluated against the stored rebuild reference.

Deliverable

A package with fixed-cell Verlet reuse enabled by explicit option and a cache statistics report.

Implemented result

Completed in 0.14.0a2. The implementation provides immutable request-keyed caches, exact S1 rebuilds at cutoff + skin, reference-relative periodic displacement checks, conservative rebuild on any cell change, exact current-frame MIC reevaluation, cache statistics, and opt-in atomic-connectivity integration. The normative specification is docs/specs/analysis/_verlet_cache_spec.md.

Stage S3 - Deformation-aware Verlet reuse

Goal

Permit safe cache reuse under variable-cell trajectories using the accepted singular-value and nonaffine-displacement bound.

Work

- compute

$$F_t = H_0^{-1} H_t;$$

- compute $\sigma_{\min}(F_t)$;
- compute correctly unwrapped nonaffine displacement by species;
- evaluate every active pair margin $M_{AB}(t)$;
- rebuild when any margin is below tolerance;
- distinguish rebuild reasons in statistics;
- handle rigid cell rotations without unnecessary rebuild;
- validate singular and ill-conditioned cells explicitly.

Acceptance gate

For every frame of every variable-cell test trajectory:

$$\text{deformation-aware cached result} = \text{fresh cell-list result} = \text{fresh dense result} .$$

Required deformation paths:

- isotropic expansion and compression;
- orthorhombic strain;

- volume-preserving shear;
- combined shear and thermal displacement;
- rigid cell rotation;
- deformation that preserves a positive margin;
- deformation that crosses the margin and triggers rebuild;
- mobile-cation and immobile-framework species with different displacement maxima;
- periodic boundary crossing during deformation;
- near-singular cells that must reject or fall back safely.

Adversarial tests should place omitted pairs just outside R_{AB} and drive the system close to the theoretical margin without crossing it.

Deliverable

A package with deformation-aware reuse available by explicit option and a mathematical validity audit.

Implemented result

Completed in 0.14.0a3. The implementation stores continuous fractional rebuild references, evaluates $F = \text{inv}(H_0)$ @ H_t , computes the smallest singular value and species-resolved nonaffine maxima, and requires every active pair margin to remain strictly above tolerance. Rigid rotations reuse safely; affine and nonaffine rebuild causes are reported separately; ill-conditioned cells are rejected explicitly. The normative specification is docs/specs/analysis/_verlet_cache_deformation_spec.md.

Stage S4 - Integration, automatic policy, benchmarks, and specifications

Goal

Integrate the verified subsystem across current consumers, establish a safe automatic backend policy, and align all documentation.

Work

- integrate persistent sessions into trajectory loops for:
 - RDF;
 - coordination distributions;
 - bond-angle distributions;
 - distance connectivity;
 - hysteretic connectivity;
 - reference connectivity;
- preserve single-frame behavior;
- expose stable high-level options, for example:

```
NeighborSearchOptions(
    backend="auto",      # auto | dense | cell_list
    cache_mode="auto",  # auto | none | verlet
    skin=0.5,
    max_consecutive_zero_reuse_rebuilds=3,
)
```

- benchmark backend crossover points;
- define auto from measured cost, not intuition;
- add diagnostics to analysis provenance;
- update specifications and architecture manuals;
- build wheel and source packages;
- run installed-wheel smoke tests.

Acceptance gate

- all legacy tests pass;
- every consumer produces identical scientific outputs under dense and optimized backends;
- full trajectory tests show exact framewise agreement;

- benchmarks demonstrate the intended scaling trend;
- auto never selects an unsupported backend;
- all Markdown and PDF specifications agree with source behavior;
- cache statistics are deterministic and auditable.

Deliverable

The first production-ready optimized neighbor subsystem.

Implemented result

Completed in 0.14.1. The public `NeighborSearchOptions` policy is integrated into RDF, coordination, bond angle, and all distance-based connectivity modes. Automatic cache resolution now requires explicit trajectory semantics: multi-frame trajectories may use deformation-aware Verlet reuse, while single-frame selections and independent ensembles remain stateless. Explicit `cache_mode="verlet"` remains available, but three consecutive completed intervals with zero successful reuse disable the cache for the remainder of the request. The deterministic backend threshold, exact fallback events, cache intervals, semantic decisions, and runtime shutoff are recorded in unified provenance. Consumer-level dense/cell-list/cache equivalence tests pass, and the normative production specification is `docs/specs/analysis/neighbor_search_spec.md`.

Test strategy

Dense equivalence is the primary oracle

Every optimized result is compared against a fresh dense calculation. Digest or state equality alone is not sufficient because two wrong geometries could produce the same graph.

The comparison must include:

- ordered center selection;
- canonical atom pair;
- periodic image shift;
- displacement vector;
- distance;
- strict cutoff inclusion;
- duplicate suppression;
- CSR offsets and grouping semantics.

Randomized property tests

Randomized tests should span:

- atom counts;
- species distributions;
- cutoff registries;
- cell lengths and skew;
- PBC masks;
- atom positions near boundaries;
- random unimodular basis transformations;
- random strains and shears;
- reference-relative displacements.

Each failure must record the random seed and full reproducing fixture.

Basis-invariance tests

For unimodularly equivalent cells

$$H' = UH,$$

with consistently transformed fractional coordinates, the original-basis physical neighbor result must be equivalent after image-shift conversion.

This directly tests reduced-basis correctness.

Cache-completeness tests

The central cache property is:

No atom pair omitted at rebuild may enter the physical cutoff while the reported safety margin remains positive.

Tests should explicitly search for counterexamples numerically near the bound.

Consumer-level regression

RDF, coordination, angle, and connectivity outputs must be compared at the final observable level, not only at the neighbor-list level.

Benchmark strategy

Benchmarks should measure both asymptotic behavior and practical crossover.

Systems

Use representative systems:

- small crystalline framework;
- replicated Na-LTA frameworks;
- dense molten salt;
- mixed zeolite-salt interface;
- highly skewed synthetic cells;
- fixed-cell and variable-cell trajectories.

Sizes

Include a geometric sequence of atom counts sufficient to expose scaling, for example:

128, 256, 512, 1024, 2048, 4096, 8192

subject to memory and runtime limits.

Metrics

Record:

- wall time per frame;
- rebuild wall time;
- reuse-frame wall time;
- peak temporary memory;
- candidate-pair count;
- accepted-pair count;
- candidate/accepted ratio;
- cache rebuild count;
- mean and median frames per rebuild;
- rebuild reasons;
- backend selected by auto;
- dense/cell-list/cached speedup.

The candidate efficiency may be summarized by

$$\eta = \frac{N_{\text{accepted}}}{N_{\text{candidate}}}.$$

A very small η indicates an oversized skin, inefficient bin layout, or an unnecessarily broad request.

Auto-backend policy

The initial auto policy should be conservative. It may use measured thresholds based on:

- estimated dense pair work;
- atom count;
- active species-pair populations;
- expected cell-list occupancy;
- whether a trajectory session can reuse a cache;
- whether the geometry satisfies supported constraints.

The policy must be deterministic and reported in provenance. Users must always be able to force `dense` or `cell_list`.

Diagnostics and provenance

The optimized subsystem should expose or record:

```
backend requested
backend selected
cache mode
skin distance
request digest
candidate pair count
accepted pair count
cell-list rebuild count
cache reuse-frame count
frames per rebuild
rebuild reason counts
minimum safety margin by rebuild interval
minimum singular value by rebuild interval
candidate/accepted ratio
fallback events
```

Diagnostics are scientific provenance and performance-tuning aids. They must not change graph identity or observable values.

Specification and documentation impact

New normative specification

Create:

```
docs/specs/analysis/neighbor_search_spec.md
docs/specs/analysis/neighbor_search_spec.pdf
```

It should own:

- dense and cell-list backend semantics;
- metric-stencil construction;
- lattice-reduction conventions;
- Verlet cache construction;
- fixed and deforming-cell validity criteria;
- request signatures;
- cache statistics;
- exactness and fallback requirements.

Existing specifications

Update the internal-neighbor specification to delegate optimized execution to the new subsystem while retaining the scientific neighbor contract.

Atomic-connectivity documentation should state that:

- geometric candidates may come from any exact backend;

- hysteretic state remains owned by atomic connectivity;
- formation and breaking classification uses one candidate-distance pass;
- cache rebuild logic is not an atomic-connectivity responsibility.

RDF, coordination, and bond-angle specifications need only brief backend-neutral statements and option documentation. Their scientific formulas and normalization must not change.

The framework/ring architecture manual should record that atomic connectivity may use the optimized backend, but no higher topology layer participates in cache management.

Risks and mitigations

Missing candidates in skewed cells

Risk: an incorrect bin stencil misses a physically close pair.

Mitigation: exact metric-box minimization, reduced-basis tests, randomized dense equivalence, and permanent dense fallback.

Wrong periodic image shift

Risk: the correct pair is found but reported in the wrong original lattice image.

Mitigation: derive vectors and shifts from one exact MIC operation, test unimodular basis transformations, and compare shifts against dense output.

Invalid cache reuse

Risk: an omitted pair enters the physical cutoff before rebuild.

Mitigation: mathematically conservative margins, adversarial near-bound tests, and framewise dense equivalence.

Wrapped-coordinate displacement errors

Risk: boundary crossings appear as large jumps or hide real motion.

Mitigation: reference-relative fractional unwrapping, explicit ambiguity detection, and boundary-crossing trajectories.

Overly large candidate lists

Risk: a large skin or skewed bin layout erases the performance benefit.

Mitigation: report candidate efficiency, reduce the lattice basis, benchmark skin choices, and retain request-specific caches.

Premature automatic selection

Risk: auto selects a slower or unsupported backend.

Mitigation: keep optimized backends explicit until benchmarks are complete; introduce auto only in S4.

Scientific-module coupling

Risk: consumers begin to own cache state or rebuild criteria.

Mitigation: centralize all geometric state in `NeighborSearchSession` and keep consumer APIs backend-neutral.

Explicitly deferred features

The first production implementation will not require:

- cross-analysis superset-cache sharing;
- pair-specific skin distances;
- GPU cell lists;
- multithreaded candidate construction;
- distributed-memory domain decomposition;
- multiple periodic images of one atom pair inside the cutoff;
- adaptive online skin optimization;

- persistent caches serialized across program runs;
- approximate or probabilistic neighbor search.

These may be added later without changing the scientific neighbor contract.

Definition of done

The complete project is done when:

1. the exact cell-list backend matches dense search for all acceptance fixtures;
2. fixed-cell cache reuse matches fresh search on every frame;
3. deformation-aware cache reuse matches fresh search on every frame;
4. hysteretic and reference connectivity use one candidate-distance pass;
5. all current scientific consumers remain backend-neutral;
6. dense, cell-list, and cached modes produce identical scientific outputs;
7. diagnostics and rebuild provenance are available;
8. benchmark results justify the automatic policy;
9. Markdown and PDF specifications match source behavior;
10. the full regression suite, build tests, and installed-wheel smoke tests pass.

Completion status

Stages S0-S4 are complete. The neighbor subsystem is production-ready under the exact single-image periodic contract. Future work may address only explicitly deferred features and must preserve the dense oracle, exact fallback, consumer neutrality, and diagnostic provenance requirements.

Appendix A - Rebuild and reuse pseudocode

```
function build_cache(frame, request):
    validate request and unique-image regime
    search_basis = reduce_lattice_if_supported(frame.cell, frame.pbc)
    bins = assign_atoms_to_fractional_bins(frame, search_basis)
    stencil = build_metric_stencil(search_basis, request.max_list_radius)
    candidates = enumerate_bin_candidates(bins, stencil)
    candidates = exact_filter_by_pair_list_radius(candidates, request)
    candidates = canonicalize_and_deduplicate(candidates)

    return VerletPairCache(
        request_digest=request.digest,
        reference_frame_index=frame.index,
        reference_positions=frame.positions,
        reference_fractional_positions=frame.fractional_positions,
        reference_cell=frame.cell,
        candidates=candidates,
        skin=request.skin,
    )

function cache_is_valid(frame, cache, request):
    if request.digest != cache.request_digest:
        return false

    F = inverse(cache.reference_cell) @ frame.cell
    sigma_min = smallest_singular_value(F)
    u_max_by_species = nonaffine_displacement_maxima(frame, cache)

    for pair_type in request.pair_types:
        margin = (
            sigma_min * (pair_type.radius + request.skin)
            - pair_type.radius
            - u_max_by_species[pair_type.A]
```

```

        - u_max_by_species[pair_type.B]
    )
    if margin <= request.safety_tolerance:
        return false

return true

function evaluate(frame, request, session):
    cache = session.cache_for(request.digest)
    if cache is absent or not cache_is_valid(frame, cache, request):
        cache = build_cache(frame, request)
        session.replace_cache(request.digest, cache)

    pair_geometry = exact_current_MIC(frame, cache.candidates)
    return apply_physical_neighbor_rule(pair_geometry, request)

```

Appendix B - Stage gate summary

Stage	New capability	Main oracle	Cache allowed?	Variable cell reuse?
S0	Dense baseline and comparison harness	Existing dense search	No	No
S1	Exact triclinic cell list	Fresh dense search	No	No
S2	Fixed-cell Verlet reuse	Fresh dense and cell list per frame	Yes	No
S3	Deformation-aware reuse	Fresh dense and cell list per frame	Yes	Yes
S4	Consumer integration and auto policy	Full observable regression	Yes	Yes

Appendix C - Accepted design decisions

The following decisions are considered settled for the initial implementation:

1. Use a reduced periodic basis when practical, but do not make correctness depend on reduction.
2. Use fractional linked cells with a metric-aware stencil, not a hard-coded 27-bin neighborhood.
3. Use pair-specific list radii and one global initial skin.
4. Retain the blocked dense backend permanently.
5. Build the cell list only on cache rebuild frames.
6. Recompute exact current MIC geometry for every cached pair on every reuse frame.
7. Use one cache per normalized request signature initially.
8. Keep hysteretic bond state separate from geometric candidate state.
9. Use the species-aware deformation margin

$$\sigma_{\min}(F_t)(r_{AB} + r_{\text{skin}}) - r_{AB} - u_A^{\max} - u_B^{\max} > \varepsilon.$$

10. Keep all rebuild decisions inside the shared neighbor subsystem.
11. Introduce automatic backend selection only after explicit-backend benchmarks.

12. Implement and release one gated stage at a time.

References

1. Quentrec, B., and Brot, C. (1973). *New Method for Searching for Neighbors in Molecular Dynamics Computations*. Journal of Computational Physics, 13(3), 430-432. DOI: 10.1016/0021-9991(73)90046-6.
2. Verlet, L. (1967). *Computer "Experiments" on Classical Fluids. I. Thermodynamical Properties of Lennard-Jones Molecules*. Physical Review, 159(1), 98-103. DOI: 10.1103/PhysRev.159.98.
3. Chialvo, A. A., and Debenedetti, P. G. (1990). *On the Use of the Verlet Neighbor List in Molecular Dynamics*. Computer Physics Communications, 60(2), 215-224. DOI: 10.1016/0010-4655(90)90007-N.
4. Heinz, T. N., and Hünenberger, P. H. (2004). *A Fast Pairlist-Construction Algorithm for Molecular Simulations under Periodic Boundary Conditions*. Journal of Computational Chemistry, 25(12), 1474-1486. DOI: 10.1002/jcc.20071.
5. Cui, Z., Sun, Y., and Qu, J. (2009). *The Neighbor List Algorithm for a Parallelepiped Box in Molecular Dynamics Simulations*. Chinese Science Bulletin, 54(9), 1463-1469. DOI: 10.1007/s11434-009-0197-0.
6. Rogers, D. M. (2016). *Overcoming the Minimum Image Constraint Using the Closest Point Search*. Journal of Molecular Graphics and Modelling, 68, 197-205. DOI: 10.1016/j.jmgm.2016.07.004.
7. Dobson, M., Fox, I., and Saracino, A. (2016). *Cell List Algorithms for Nonequilibrium Molecular Dynamics*. Journal of Computational Physics, 315, 211-220. DOI: 10.1016/j.jcp.2016.03.056.
8. Larsen, A. H., et al. (2017). *The Atomic Simulation Environment - A Python Library for Working with Atoms*. Journal of Physics: Condensed Matter, 29(27), 273002. DOI: 10.1088/1361-648X/aa680e.
9. Nguyen, P. Q., and Stehlé, D. (2009). *Low-Dimensional Lattice Basis Reduction Revisited*. ACM Transactions on Algorithms, 5(4), Article 46. DOI: 10.1145/1597036.1597050.