

MLFF Architecture Revision 53

2026-08-15

Revision 53

Revision 53 implements **VRAM1 + PERF-P4** without interrupting development for GPU qualification. Accelerator-memory and throughput acceptance remains scheduled under FINAL -GPU1; **PERF-P5** becomes the next CPU/control-plane implementation gate.

Capacity evidence

MaceBatchCapacityCalibration.v2 now measures the actual descriptor, prediction, or combined DATA6 workload. It binds deterministic stress-frame identities, per-probe throughput, CUDA allocated/reserved peaks, driver free/total memory, absolute/fractional headroom, host-residency estimates, and post-cleanup state. Historical v1 remains readable with descriptor-only semantics.

A calibrated cap is only a prior. DATA6 re-clamps it against fresh live VRAM before execution. Runtime OOM backoff is durable and identity-bound, so a matching restart reuses the learned safe cap while changed model/policy/device/workload/calibration identity invalidates it.

Bounded execution pipeline

Native MACE combined evaluation can prebuild batch $n + 1$ on CPU while batch n is evaluated, and shard persistence can overlap behind both. Queues are bounded and included in host-memory admission. Persistence completion never changes scientific order; the append-only journal and ordered compaction remain recovery authority. Synchronous execution is retained as an exact fallback.

Pinned/nonblocking transfers are not promoted without accelerator evidence. This follows PyTorch's documented workload-dependent transfer behavior and preserves the project rule that execution optimizations must be measured rather than assumed.

Reference qualification

The supplied MH-1/omat_pbe and MPA-0-medium/default CPU/e3nn prepared and direct combined-evaluation paths are byte-identical on locked regression structures. The bounded 44-frame DATA6 fixture also yields one identical scientific signature under synchronous and pipelined orchestration. The pipeline is 4.72% slower on that tiny CPU-only fixture, so no CPU speedup claim is made.

Qualification boundary

The development host cannot qualify CUDA peak/reserved memory, OOM capacity, transfer overlap, or accelerator throughput. Those items remain explicit FINAL -GPU1 obligations. This release therefore records VRAM1/PERF-P4 as **implemented and CPU/control-plane qualified, accelerator qualification pending**.

Next implementation gate

PERF-P5 is next. Its CPU/control-plane implementation may proceed before the deferred E3NN-BASELINE campaign; final performance qualification remains conditional on the release-matched GPU baseline.